

Machine Learning

Artificial Neural Networks







mite

container ship

motor scooter

leopard

	mite	container ship	motor scooter	leopard
	black widow	lifeboat	go-kart	jaguar
	cockroach	amphibian	moped	cheetah
	tick	fireboat	bumper car	snow leopard
	starfish	drilling platform	golfcart	Egyptian cat

			
grille	mushroom	cherry	Madagascar cat
convertible	agaric	dalmatian	squirrel monkey
grille	mushroom	grape	spider monkey
pickup	jelly fungus	elderberry	titi
beach wagon	gill fungus	ffordshire bullterrier	indri
fire engine	dead-man's-fingers	currant	howler monkey

Overview:


- Artificial Neuron
- Activation Function
- Neuron Capacity
- Single Layer ANN
- Multilayer ANN
- Universal Approximator
- Motivation behind ANN
- Backpropagation

- Hidden unit pre-activation:

$$z(x) = \sum_i w_i x_i + b = W^T X + b$$

- Hidden unit activation:

$$f(x) = \sigma(z(x)) = \sigma(\sum_i w_i x_i + b)$$



\mathbf{W} are weight matrix connects i^{th} hidden unit with i^{th} input unit


\mathbf{b} are bias vectors

$\sigma(\cdot)$ is the activation function

Sigmoid:

$$\sigma(z) = \text{sigm}(z) = \frac{1}{1 + \exp(-z)}$$


- Squashes the hidden unit's pre-activation to between 0 and 1.
- Always positive.
- Bounded.
- Strictly increasing.



Hyperbolic tangent (“tanh”):

$$\sigma(z) = \tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)} = \frac{\exp(2z) - 1}{\exp(2z) + 1}$$

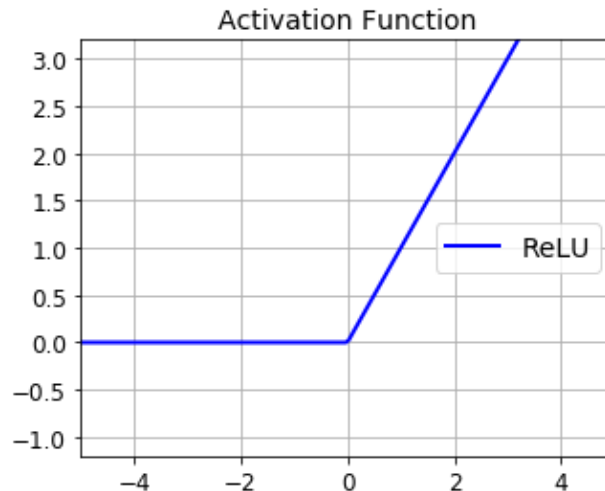
- Squashes the hidden unit's pre-activation to between -1 and 1.
- Can be positive or negative.
- Bounded.
- Strictly increasing.



Rectified Linear Unit (“ReLU”):

$$\sigma(z) = \text{ReLU}(z) = \max(0, z)$$

- Bounded below by 0 (always non-negative).
- Not bounded above.
- Strictly increasing.



- Hidden layer pre-activation:

$$z(x)_i = b_i^{(1)} + W_{i,j}^{(1)} x_j$$

Similarly in Matrix form:


$$z(X) = b^{(1)} + W^{(1)} X$$

- Hidden layer activation:

$$h(X) = \sigma(z(X))$$

- Output layer activation “ Φ ”:

$$F(X) = \Phi(b^{(2)} + W^{(2)T} h^{(1)}X)$$



SoftMax Activation Function

- Multi-class classification:
 - requires multiple outputs i.e. 1 output per class.
 - need to estimate the conditional probability of output belonging to a particular class c , $p(y = c | \mathbf{x})$.

- Apply the SoftMax activation function at the output:

$$\Phi(z) = \text{SoftMax}(z) = \left[\frac{e^{z_1}}{\sum_c e^{z_c}}, \dots, \frac{e^{z_1}}{\sum_c e^{z_c}} \right]^T$$

- strictly positive
- sums to one
- Predicted class is the one with highest estimated probability

Multilayer NN with L hidden layers

- **Hidden layer pre-activation for $k > 0$:**

$$h^{(0)}(X) = X,$$


$$z^{(k)}(X) = b^{(k)} + \mathbf{W}^{(k)} h^{(k-1)}(X)$$

- **Hidden layer activation (k from 1 to L):**

$$h^{(k)}(X) = \sigma(z^{(k)}(X))$$


- **Output layer activation ($k = L+1$):**

$$F(X) = h^{(L+1)}(X) = \Phi(z^{(L+1)}(X))$$



Capacity of Single Hidden Unit

- Range of hidden unit determined by $\sigma(\cdot)$
- Bias b changes the position of the riff.





Bias only changes
the position of the
riff

Source: Pascal Vincent


Capacity of Single Hidden Unit

- Could do binary classification.
- With sigmoid, can interpret neuron as estimating $p(y = 1 | \mathbf{x})$.
- Also known as logistic regression classifier, if greater than 0.5 predict class 1, otherwise predict class 0.





Capacity of Single Hidden Unit

- Can solve linearly separable problems



Capacity of Single Hidden Unit

- Cannot solve non-linearly separable problems
- Unless the input is transformed into a separable representation



- Hidden layer pre-activation:

$$z(x)_i = b_i^{(1)} + W_{i,j}^{(1)} x_j$$

Similarly in Matrix form:


$$z(X) = b^{(1)} + W^{(1)} X$$

- Hidden layer activation:


$$h(X) = \sigma(z(X))$$

- Output layer activation “ Φ ”:

$$F(X) = \Phi(b^{(2)} + W^{(2)T} h^{(1)} X)$$




Capacity of Single Hidden Layer Neural Network




Source: Pascal Vincent

Capacity of Single Hidden Layer Neural Network



Source: Pascal Vincent

Capacity of Single Hidden Layer Neural Network



- **Universal Approximation Theorem (Hornik, 1991)**

“Single layer feedforward network can approximate any continuous function arbitrarily well if and only if the network’s activation function is continuous, non-constant, and bounded.”

- Hornik’s result applies for sigmoid, tanh and many other hidden layer activation functions.
- However, modern day defacto activation function is ReLu, and it does not satisfy Hornik’s theorem as the $\text{ReLu}(z) = \max(0, z)$ is unbounded from above.

- “Multilayer feedforward network can approximate any continuous function arbitrarily well if and only if the network's continuous activation function is not polynomial.”

Definition

A set F of functions in $L_{loc}^{\infty}(R^n)$ is dense in $C(R^n)$ if for every function $g \in C(R^n)$ and for every compact set $K \subset R^n$, there exists a sequence of functions $f_j \in F$ such that

$$\lim_{j \rightarrow \infty} \|g - f_j\|_{L^{\infty}(K)} = 0.$$

Theorem

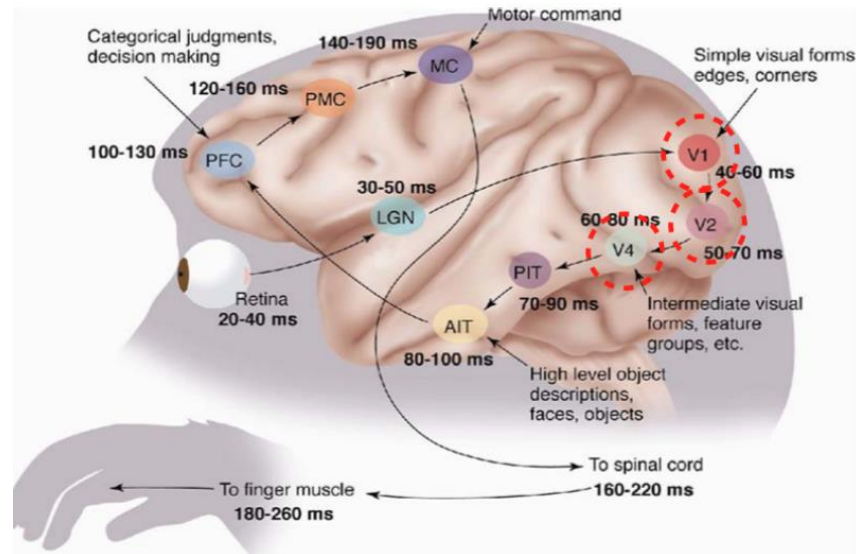
(Leshno et al., 1993) Let $\sigma \in M$, where M denotes the set of functions which are in $L_{loc}^{\infty}(\Omega)$.

$$\Sigma_n = \text{span}\{\sigma(w \cdot x + b) : w \in R^n, b \in R\}$$

Then Σ_n is dense in $C(R^n)$ if and only if σ is not an algebraic polynomial (a.e.).

Biological Inspiration


- Parallel with the visual vortex



Source: Simon Thorpe


Biological Inspiration

- Parallel with the visual vortex



Source: Simon Thorpe



- Parallel with the visual vortex



Source: Simon Thorpe

Biological Inspiration


- Parallel with the visual vortex
- Edges



Source: Simon Thorpe

Biological Inspiration


- Parallel with the visual vortex
- Edges
- Higher level features such as nose, mouth, eyes
- Face



- **Firing rates of different input neurons combine to influence the firing rate of other neurons:**
 - depending on the dendrite and axon, a neuron can either work to increase (excite) or decrease (inhibit) the firing rate of another neuron
- **This is what artificial neurons approximate:**
 - the activation corresponds to a “sort of” firing rate
 - the weights between neurons model whether neurons excite or inhibit each other
 - the activation function and bias model the threshold behavior of action potentials

Forward propagation

- Randomly Initialize Θ
- $\Theta = \{W^{(1)}, b^{(1)}, W^{(2)}, b^{(2)}, \dots, W^{(L+1)}, b^{(L+1)}\}$




- Objective Function for Multi-class Classification

$$\underset{\boldsymbol{\theta}}{\operatorname{argmin}} \quad \frac{1}{m} \sum_{i=1}^m l(\mathbf{F}(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}) + \lambda \boldsymbol{\Omega}(\boldsymbol{\theta})$$

- Loss Function: Negative log likelihood


$$l(\mathbf{F}(\mathbf{x}, \boldsymbol{\theta}), \mathbf{y}) = -\sum_c \mathbf{1}_{(y=c)} \log \mathbf{F}(\mathbf{x})_c$$



Cho & Chow, Neurocomputing 1999

Backpropagation

- Optimization Algorithm e.g. stochastic gradient descent
- Objective Function
- Gradient of Output Layer
- Gradient of Hidden Layer
- Gradient of Activation Function
- Gradient of Parameters



Optimization Algorithm

Stochastic (incremental) gradient descent


For each training example ($\mathbf{x}^{(i)}, \mathbf{y}^{(i)}$)

For N iterations

- Compute $\Delta = - \nabla_{\boldsymbol{\theta}} l(\mathbf{F}(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)}) - \nabla_{\boldsymbol{\theta}} \lambda \Omega(\boldsymbol{\theta})$
- Update $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha \Delta$


Loss Gradient at Output Layer

- $$\nabla_{F(\mathbf{x})} - \log F(\mathbf{x})_y = \frac{-e(y)}{F(\mathbf{x})_y}$$



Loss Gradient at Output Pre-activation


- $\nabla_{\mathbf{z}^{(L+1)}(\mathbf{x})} - \log \mathbf{F}(\mathbf{x})_y = -(\mathbf{e}(\mathbf{y}) - \mathbf{F}(\mathbf{x}))$



Loss Gradient at Hidden Layer


- $$\nabla_{\mathbf{z}^{(k)}(\mathbf{x})} - \log \mathbf{F}(\mathbf{x})_y$$

$$= \nabla_{\mathbf{h}^{(k)}(\mathbf{x})} - \log \mathbf{F}(\mathbf{x})_y \odot [\dots, \sigma'(\mathbf{z}^{(k)}(\mathbf{x})) , \dots]$$




Sigmoid Function Derivative:

- $\sigma(z) = \text{sigm}(z) = \frac{1}{1+\exp(-z)}$
- $\sigma'(z) = \sigma(z) (1 - \sigma(z))$



Hyperbolic tangent (“tanh”) Derivative:


- $\sigma(z) = \tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)} = \frac{\exp(2z) - 1}{\exp(2z) + 1}$
- $\sigma'(z) = 1 - \sigma(z)^2$



Gradient of Weights

- $$\nabla_{\mathbf{w}^{(k)}} - \log \mathbf{F}(\mathbf{x})_y$$


$$= (\nabla_{\mathbf{a}^{(k)}(\mathbf{x})} - \log \mathbf{F}(\mathbf{x})_y) \mathbf{h}^{(k-1)}(\mathbf{x})^\top$$



Gradient of Biases

- $$\nabla_{\mathbf{b}^{(k)}} - \log \mathbf{F}(\mathbf{x})_y$$

$$= (\nabla_{\mathbf{a}^{(k)}(\mathbf{x})} - \log \mathbf{F}(\mathbf{x})_y)$$



Books

- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Aurélien Géron
- Pattern recognition and machine learning, Christopher M. Bishop
- The Elements of Statistical Learning, Hastie, Tibshirani, and Friedman

Blogs, code snippets, lecture notes, etc.

- <https://github.com/stephencwelch/Neural-Networks-Demystified> This tutorial uses numpy and python
- <http://www.wildml.com/2015/09/implementing-a-neural-network-from-scratch> This tutorial implements a classifier in python from scratch
- <http://deeplearning.net/tutorial> This set of tutorials uses the Theano package, which is pretty tricky to learn (but very powerful because it calculates gradients for you automatically, among other things)
- <http://www.cs.stir.ac.uk/courses/ITNP4B/lectures/kms/1-Intro.pdf> For some of the history and possible connections to neuroscience
- http://www.dmi.usherb.ca/~larocheh/index_en.html Modified from Hugo Larochelle